

The Data Moat Thesis

Why Compounding Data Is the Most Defensible Asset Class— and Why It Follows the Physics of Black Holes

Ryan Barrett

<https://github.com/elevate-foundry>

May 2026

*Lat. **accrēscere**, to grow toward—from the same root as **creāre**, to bring forth.
Creation is accretion.*

Abstract. It takes 426 bytes to describe your entire financial relationship with a creditor. That single Metro 2 record—balance, payment history, account status—is a snapshot. It has value for exactly as long as it is current, and then it decays. But layer 120 of those snapshots across 30 creditors over 7 years, and you have something qualitatively different: a *trajectory*. A pattern. A predictive surface that no competitor can fabricate because it required years to accumulate.

This paper argues that **data accretion**—the continuous, consent-bounded gathering, cleaning, connecting, and enriching of scattered data over time—creates the most defensible competitive asset in the information economy. We propose that accreting data systems follow the same mechanics as astrophysical black holes: there is an event horizon (consent boundary), an accretion disk (ingestion pipeline), gravitational compression (memory core), and—critically—an information density beyond which the accumulated dataset becomes irreplaceable. We call this the **data singularity threshold**: the point at which a historical dataset is so rich, so interconnected, and so temporally deep that it constitutes a permanent structural advantage.

This is not metaphor. It is mechanics.

1. The 426-Byte Seed

The Metro 2 format is the universal language of American consumer credit. Maintained by the Consumer Data Industry Association (CDIA), it defines exactly how creditors report account information to the three major bureaus. Each tradeline segment is precisely **426 bytes**—a fixed-width record containing:

- Account type and status (2 bytes)
- Payment rating history (24 bytes—two years of monthly snapshots)
- Current balance, credit limit, scheduled payment (10 bytes each)
- Date opened, last payment, last activity (6 bytes each)
- Consumer identifiers—SSN, name, DOB, address (~200 bytes)

A complete credit file—typically 30 to 50 tradelines—runs approximately 13,000 to 21,000 bytes. **Twenty**

kilobytes to describe an entire financial life as the credit system knows it. That is smaller than the average JPEG thumbnail.

And yet this 20 KB file determines whether a consumer gets the apartment, the car, the house, the business loan. It determines insurance premiums. It determines whether a landlord returns a call. The *density* of consequential information per byte is extraordinary.

But here is the critical property of Metro 2: **each record is a snapshot**. It tells you what is true *right now*. The payment history field gives 24 months of lookback, and then the past disappears. Seven years after an account closes, it vanishes from the file entirely.

Snapshots decay. That is their nature. What if they did not?

2. Snapshots vs. Accretion

The fundamental distinction in data strategy is between **snapshot systems** and **accretion systems**.

A *snapshot system* captures the current state. A CRM says the lead works at Acme Corp with a \$50K budget. An analytics dashboard shows 4,200 monthly active users. A credit report shows a 740 FICO score. Each is true at the moment of observation and begins degrading immediately.

An *accretion system* captures the current state *and* preserves every prior state, connecting them into a continuous historical record. The lead *used to* work at Initech, moved to Acme Corp eight months ago after a Series B (known because a press release was also ingested), and their budget doubled since Q2 (known because Q3 earnings were ingested). The MAU is 4,200, but the system knows it was 3,100 six months ago, 1,800 a year ago, and can project the trend. The FICO is 740, but the trajectory shows recovery from 580 after a bankruptcy discharge three years prior—a very different risk profile than a steady-state 740.

The snapshot says *what*. The accretion says *what, when, why, and what comes next*.

2.1 The Compounding Effect

Accreted data compounds. Each new data point does not merely add to the collection—it *enriches every prior data point* by providing context, confirming patterns, or revealing contradictions. This is not linear growth. It is network-effect growth applied to information:

Day	Signals	Value
1	Name, email	Minimal
30	+ company, role, 3 web visits	Moderate
180	+ firmographics, tickets, webinar	Significant
365	+ behavior patterns, job change, NPS	Compound

Table 1: Progressive enrichment of a single entity profile through continuous data accretion.

The Day 365 dataset is not 12× more valuable than the Day 30 dataset. It is *categorically different*. It enables predictions that were impossible

at any prior stage. And—critically—it cannot be replicated by a competitor who starts their accretion process today. They are structurally 365 days behind, and that gap widens with every day of continued accretion.

This is the data moat.

2.2 Worldlines: The Causal Primitive

The distinction between snapshot and accretion is not merely temporal—it is *causal*. In physics, a **worldline** is the path an object traces through spacetime: not just where it is, but where it has been, and the causal ordering of its states.

$$\text{Worldline}(\text{entity}) = \{(t_0, S_0), (t_1, S_1), \dots, (t_n, S_n) \mid t_0 < t_1 < \dots\} \quad (1)$$

where each (t_i, S_i) is a state at time t_i , and the sequence is ordered by causality.

A snapshot system stores:

$$\text{Snapshot} = S_n \quad (\text{state at } t_{\text{now}}) \quad (2)$$

An accretion system stores:

$$\text{Accretion} = \text{Worldline} \quad (\text{full trajectory from } t_0 \text{ to } t_n) \quad (3)$$

The critical insight: **two identical snapshots can arise from different worldlines**. A FICO score of 740 today could result from:

- Stable trajectory: 740 for the past 5 years
- Recovery trajectory: 580 three years ago, recovering to 740 today
- Deterioration trajectory: 800 three years ago, declining to 740 today

Each worldline implies different risk. Only the worldline—the causal history—distinguishes them. Therefore:

$$\text{Risk} = f(\text{Worldline}), \quad \text{not} \quad f(\text{Snapshot}) \quad (4)$$

Your system does not store entities. It stores *worldlines*. Each entity is a particle moving through a space of financial and behavioral states. The accretion disk constructs these worldlines by ingesting

state transitions over time. The event horizon determines which worldlines enter the system (consent boundary). The memory core compresses worldlines into four tiers (Raw $\rightarrow$ Normalized $\rightarrow$ Semantic $\rightarrow$ Symbolic), each a progressive compression of the full trajectory. Hawking radiation extracts *projections* of worldlines: FICO is a scalar projection; Metro 2 is a structured projection; your system returns the worldline itself.

The singularity threshold is the point where the worldline is so deep and interconnected that its causal structure is irreducible. A competitor cannot fabricate causal history. They can only begin accreting today, and they will always be behind on worldline depth.

3. The Event Horizon Made Tactile

The event horizon—the consent boundary—is the most important structural element of the data accretion system. Without it, the system has no legal standing, no trust, and no moat. But an event horizon that only sighted, literate, English-speaking users can interact with is not a universal boundary. It is a partial one.

How does a blind person verify that their data is in the system? How do they read their own consent scope? How do they query their own worldline?

If they cannot, the event horizon excludes them—and the system’s effective mass (N) shrinks. Accessibility is not a feature appended to the architecture. It is the *physical implementation* of the consent boundary.

3.1 Identity Encoding as Consent Boundary

The jurisdiction-configurable identifier set $\mathcal{S}_{\text{national_id}}$ solves the classification problem but not the *representation* problem. An SSN is 9 decimal digits (30 bits). An Aadhaar is 12 digits (40 bits). A BSN is 9 digits (30 bits). A Steuer-ID is 11 digits (37 bits). Each has a different length, format, check-digit algorithm, and regulatory context. At the data layer, this heterogeneity creates fragility.

We propose an n -dot binary encoding inspired by Braille [12]. Standard Braille uses a 2×3 matrix of raised dots, yielding $2^6 = 64$ unique cells. Computer

Braille extends this to 2×4 ($2^8 = 256$ patterns)—one full byte. The generalization:

$$\text{Capacity}(n) = 2^n \text{ patterns} \quad (5)$$

n	Patterns	Can Encode	Layout
6	64	Single character set	2×3
8	256	Full ASCII / one byte	2×4
10	1,024	Unicode BMP block	2×5
30	~ 1 B	SSN, BSN (9 decimal digits)	2×15
40	~ 1 T	Aadhaar (12 decimal digits)	2×20
128	3.4×10^{38}	UUID / any global identifier	2×64

Table 2: n -dot encoding capacity and representative uses.

Any national identifier can be represented as an n -dot binary cell, where $n = \lceil \log_2(\text{max_value}) \rceil$. This is not a clever encoding trick. It is the event horizon made tangible:

1. **Uniform storage:** Every identifier is a fixed-length bit vector, regardless of jurisdiction
2. **Tactile verification:** The $2 \times \lceil n/2 \rceil$ matrix is physically producible on refreshable Braille displays—an entity can *touch* their own identity
3. **Hashing compatibility:** n -dot representations compose naturally with HMAC-based tokenization
4. **Jurisdiction abstraction:** The policy engine operates on bit vectors; jurisdiction-specific validation is isolated to the ingestion boundary

3.2 The Encoding Function

Given a national identifier id from jurisdiction j :

$$\text{BRAILLE_ENCODE}(id, j) = \langle \text{prefix}_j \parallel \text{bin}(id) \parallel \text{check}(id, j) \rangle_{n_j} \quad (6)$$

where prefix_j is a jurisdiction tag (e.g., 3 bits for country code), $\text{bin}(id)$ is the binary representation, check is the jurisdiction-specific check digit recomputed in binary, and n_j is the total dot count. A 30-bit SSN becomes a 2×17 Braille cell (30 bits + 3 prefix + 1 check). An Aadhaar becomes 2×22 .

This encoding is not theoretical. Refreshable Braille displays (e.g., HumanWare Brailliant, HIMS) al-

ready render 8-dot cells at 40–80 cells per line. An n -dot identifier occupies 2–5 cell positions on existing hardware.

3.3 Consent Made Tangible

Each entity’s worldline (Section 2) is encrypted under a per-entity key. That key is represented in n -dot Braille. The entity can:

1. **Verify their identity** in Braille on a refreshable display
2. **Read consent artifacts** in ARIA-compatible formats (screen-reader, structured JSON, audio summary)
3. **Query the system** via voice-initiated retrieval
4. **Receive results** in haptic, audio, and Braille—non-visual alternatives to every visualization
5. **Revoke consent** through the same accessible interface, triggering crypto-shredding of their per-entity key

This is the event horizon made physical. Consent is not a checkbox buried in a Terms of Service. It is a boundary that the entity can *touch*, *hear*, and *verify*.

3.4 The Accessibility Stack

1. **Identity layer:** n -dot Braille encoding ensures identifiers are natively representable on tactile hardware
2. **Consent layer:** Consent artifacts in screen-reader-compatible formats (ARIA-annotated HTML, structured JSON, audio summary)
3. **Audit layer:** Every policy decision (ALLOW, REDACT, QUARANTINE, BLOCK, DELETE) produces human-readable *and* machine-parseable output consumable by assistive technology
4. **Query layer:** Purpose-constrained retrieval supports voice-initiated queries and Braille-rendered results
5. **Visualization layer:** Timeline visualizations provide non-visual alternatives (sonification, haptic feedback, structured text)

3.5 Why the Moat Requires Accessibility

Accessibility is not orthogonal to competitive advantage—it *amplifies* it:

- **Larger consent surface:** An accessible system can accrete data from populations that inaccessible systems exclude. This directly increases N in the effective information equation ($I_{\text{eff}} = N \times S$)
- **Regulatory durability:** Accessibility compliance (ADA, Section 508, EN 301 549, WCAG 2.2) reduces legal attack surface and strengthens the consent boundary
- **Trust signal:** Demonstrable accessibility builds trust with regulators, advocacy organizations, and the members themselves—accelerating consent conversion

Credit unions serve disproportionate numbers of members with disabilities, elderly populations, and underbanked communities. A system that cannot serve them has a smaller event horizon—and a smaller moat. The event horizon is only as strong as its weakest point of interaction.

4. The Physics of Information Gravity

The black hole metaphor is not decorative. Data accretion follows the same mechanics as astrophysical accretion, and the mapping is precise enough to be instructive.

4.1 The Event Horizon: Consent as Gravity’s Boundary

In astrophysics, the event horizon is the boundary beyond which nothing escapes. In data accretion, the **consent boundary** serves an identical function: data can enter the system only through authorized, consented channels.

A black hole that absorbed everything indiscriminately would be incoherent—a soup of unrelated matter with no usable structure. Similarly, a data system that ingests without consent, classification, and policy gates accumulates noise, liability, and legal risk rather than value.

The consent boundary is what makes accreted data *trustworthy*. Every record has provenance: where it came from, who authorized it, what purposes it serves, when consent was granted, when it expires. This provenance is itself a compounding asset.

4.2 The Accretion Disk: Pipeline as Compression Engine

Around a black hole, infalling matter does not plunge directly in. It spirals through an accretion disk—a region of intense heating, compression, and radiation where matter is transformed before crossing the event horizon.

The data ingestion pipeline is this disk. Raw input—CSV, webhook payload, web page, pasted paragraph—spirals through stages of increasing refinement:

1. **Parsing:** Raw bytes → structured fields
2. **Classification:** Sensitive data identified and labeled
3. **Policy evaluation:** ALLOW, REDACT, QUARANTINE, BLOCK, DELETE
4. **Normalization:** Heterogeneous formats → canonical form
5. **Entity resolution:** Separate records → same real-world entity
6. **Compression:** Full text → summary → keywords → symbolic
7. **Embedding:** Semantic meaning encoded as vector

By the time data crosses the event horizon into the memory core, it has been transformed from raw material into structured, classified, deduplicated, consented, compressed knowledge. The pipeline *adds value* to every byte—just as the accretion disk converts gravitational potential energy into radiation.

4.3 The Singularity: Compressed Knowledge Core

At the center of a black hole, all accreted matter is compressed into a singularity. The memory core serves the same function via a four-tier compression hierarchy:

- **Tier 1—Raw:** Original content, post-redaction
- **Tier 2—Normalized:** Structured, standardized canonical form
- **Tier 3—Semantic:** Summaries, extracted relationships, key phrases
- **Tier 4—Symbolic:** Keywords, entity links, category tags

Each tier trades fidelity for density. Together, they provide flexibility: deep retrieval for precision, fast pattern matching for speed.

4.4 Hawking Radiation: Value Escaping the System

Hawking’s most famous insight was that black holes radiate. Information escapes, slowly, as thermal radiation.

In an accretion system, **retrieval is Hawking radiation**: value escaping the compressed core as answers, predictions, alerts, and exports. The output is proportional to the mass: a larger, denser memory core produces more valuable, more accurate, more contextual radiation.

4.5 Time Dilation: Why History Outweighs Currency

Near a black hole, time slows. A clock at the event horizon ticks slower than a distant clock.

In data accretion, historical depth creates an analogous effect: **older data acquires disproportionate value when contextualized by newer data**. A job title from 2022 is stale in isolation. In the context of three subsequent job changes, a funding event, and a product launch, it reveals a career trajectory that predicts future behavior.

This is the fundamental asymmetry: *time cannot be compressed*. A competitor with superior algorithms and deeper pockets still cannot replicate two years of continuously accreted data. They can begin today, but they will always be two years behind on temporal depth.

5. The Physics of the Threshold

Before we map these concepts to markets, it is worth quantifying the information-theoretic limits that make the metaphor precise.

5.1 Landauer’s Principle

The minimum energy required to erase (or equivalently, to store) one bit of information at temperature T is [1]:

$$E_{\min} = k_B T \ln 2 \quad (7)$$

where $k_B = 1.381 \times 10^{-23}$ J/K is the Boltzmann constant. At room temperature ($T = 300$ K):

$$E_{\text{bit}} \approx 2.87 \times 10^{-21} \text{ J} \quad (8)$$

A single Metro 2 record (426 bytes = 3,408 bits) therefore carries a minimum thermodynamic cost of:

$$E_{\text{Metro 2}} = 3,408 \times 2.87 \times 10^{-21} \approx 9.8 \times 10^{-18} \text{ J} \quad (9)$$

5.2 The Bekenstein Bound

The Bekenstein bound [2] establishes the maximum entropy (information) that can be contained within a spherical region:

$$S_{\max} = \frac{2\pi R E}{\hbar c \ln 2} \quad (10)$$

where R is the radius, E is the total energy, $\hbar$ is the reduced Planck constant, and c is the speed of light. When the information content of a region reaches this bound, the region *must* collapse into a black hole.

5.3 Metro 2 Records to Black Hole

To form the smallest possible black hole (Planck mass, $m_P = 2.176 \times 10^{-8}$ kg), the required energy is:

$$E_P = m_P c^2 = 1.956 \times 10^9 \text{ J} \quad (11)$$

At Landauer-minimum energy per byte ($E_{\text{byte}} = 2.30 \times 10^{-20}$ J):

$$N_{\text{bytes}} = \frac{E_P}{E_{\text{byte}}} = \frac{1.956 \times 10^9}{2.30 \times 10^{-20}} \approx 8.5 \times 10^{28} \text{ bytes} \quad (12)$$

In Metro 2 records:

$$N_{\text{Metro 2}} = \frac{8.5 \times 10^{28}}{426} \approx 2.0 \times 10^{26} \text{ tradelines} \quad (13)$$

Two hundred septillion credit tradelines—approximately one thousand tradelines for every star in the observable universe—compressed to Landauer-minimum energy density, would form a gravitational singularity.

The number is absurd in the literal sense. But the *structure* of the argument is not. Information has physical weight. Accumulation has gravitational consequence. And there exists a density threshold beyond which the accumulated information becomes irreducible.

In data strategy, the singularity threshold is not physical—it is economic. It is the point at which accreted data becomes so interconnected and temporally deep that replication is structurally infeasible. The physics provides the framework; the market determines the scale.

6. The Four Value Domains

Data accretion creates defensible value across four primary domains. Each follows the same compounding dynamic.

6.1 Knowledge & Intelligence Gathering

Thesis: Progressively building profiles or knowledge bases that deepen with every new data point.

A firm monitoring a competitor’s job postings, pricing, and press releases has a snapshot on day one. After six months, it detects hiring surges (expansion signal), pricing changes (market pressure), and stack shifts (strategic direction). After two years, it has a predictive model correlating hiring patterns with product launches at 80%+ accuracy.

Similarly, OSINT investigations layer social media, corporate filings, property records, and news mentions. Each source alone reveals fragments. The accretion reveals connections: the LLC sharing a registered agent with two others, all transacting with the same shell company during the same quarter as an unusual SEC filing. These patterns require temporal depth.

Moat: Historical pattern depth—the ability to say “this has happened before, and here is what followed.”

6.2 Analytics & Recommendation Engines

Thesis: Accreted behavioral data powers personalization that improves with every interaction.

Netflix’s recommendation engine is not valuable because of its algorithm—the algorithm is published. It is valuable because it sits atop billions of viewing histories accumulated over decades. A new entrant with the same algorithm but zero history will produce inferior recommendations until comparable behavioral depth is accreted.

A business connecting Stripe, Shopify, CRM, support, and marketing into an accretion engine initially reports revenue by channel. After a year, it correlates CAC with LTV with support burden with churn probability—per cohort, per channel, per product, with seasonal adjustment.

Moat: Behavioral depth and cross-source temporal correlation.

6.3 AI/ML Training Data Curation

Thesis: Model quality is bounded by training data quality, and continuously curated data is a compounding asset.

Every frontier model is trained on web-scale data. But enterprise AI differentiation is not web-scale breadth—it is **domain-specific depth**. A firm accreting cleaned, labeled, consent-bounded transaction data for three years has a corpus that is higher quality than web scrapes, domain-specific, temporally structured, fully provenanced, and continuously growing.

A competitor cannot download this corpus. They cannot buy it. They cannot generate it synthetically. They can only begin their own accretion and wait.

Moat: Curated, consent-bounded, domain-specific training corpora.

6.4 Immutable Histories & Compliance

Thesis: Continuous, auditable records become legally required proof.

Institutions under SOX, Basel III, or GDPR must demonstrate what data they held, when, what decisions they made, and why. An append-only audit

log tracking every ingestion, decision, retrieval, and deletion is not a feature—it is a license to operate.

A manufacturer accreting shipment logs, vendor data, and quality inspections can answer: “When did we know Supplier X’s defect rate exceeded threshold, and what did we do?” That question is asked by regulators, insurers, and litigators. The answer exists only in a system that was continuously accreting at the time.

Moat: Immutable temporal proof—demonstrating what was known and when.

7. Hidden Challenges

Data accretion is not “save everything forever.” Three challenges must be architecturally resolved.

7.1 Data Staleness and Degradation

Data decays. A job title from 2022 may be wrong today. Pricing from six months ago is likely stale. Naive accretion—storing everything with equal weight—produces a system where stale data pollutes fresh data, degrading predictions and eroding trust.

Required mechanisms:

- *Staleness scoring:* Time-decay function reduces retrieval weight
- *Chronological weighting:* Fresh data ranks higher by default
- *Automated deprecation:* Records exceeding thresholds are flagged or archived
- *Re-verification:* Source refreshes reset staleness
- *“As-of” queries:* Retrieve the knowledge state at any historical date

The key insight: **staleness management is not deletion**. Stale data retains historical value for trajectory analysis. It should lose *retrieval priority* without losing *existence*.

7.2 Ethical and Legal Ramifications

GDPR (Europe) and CCPA/CPRA (California) establish data subject rights that directly conflict with unbounded accretion:

- **Right to be forgotten:** The system must identify and purge every record containing a sub-

ject’s information—across all four memory tiers, including vector embeddings

- **Purpose limitation:** Data collected for one purpose cannot be repurposed without consent; retrieval must be purpose-constrained and auditable
- **Data minimization:** Collection must not exceed what is necessary for stated purposes

The resolution is architectural: **consent is not a checkbox; it is a runtime enforcement layer.** Every record carries its consent scope. Every query declares its purpose. The system evaluates access dynamically. The event horizon is what makes accretion legally viable.

7.3 Database Architecture

Standard relational databases buckle under continuous ingestion of complex, interconnected, temporal data. Three storage paradigms are required:

- **Relational** (PostgreSQL): Transactional data, source metadata, policy rules
- **Vector** (pgvector): Semantic search over embedded documents; HNSW/IVFFlat indexing for sub-second retrieval across millions of vectors
- **Time-series** (TimescaleDB): Temporal queries—“How did Entity X change between January and June?”—optimized for sequential, timestamped data

The discipline is to use each engine for its designed query pattern rather than forcing one to serve all three.

8. The Gravity Equation

We return to the physics. The gravitational force between two masses is:

$$F = G \frac{m_1 \cdot m_2}{r^2} \quad (14)$$

In data accretion, this maps directly:

Physics	Data Strategy
m_1	Size and depth of accreted dataset
m_2	Importance of the decision being informed
r	Freshness gap (recency of relevant data)
F	Pull of the dataset—ability to attract users and resist competitive displacement
G	Quality of the accretion architecture

Table 3: Mapping Newton’s gravitational force to data accretion dynamics.

As m_1 grows (more data accreted over more time), F increases. As r shrinks (fresher data), F increases further. The force is proportional to the product of depth and decision importance, inversely proportional to the square of the staleness gap.

This is why accreted data is the most defensible asset class:

1. **It compounds:** m_1 grows with every ingestion
2. **It requires time:** Temporal depth cannot be compressed
3. **It resists replication:** A competitor starting today is structurally behind
4. **It improves with use:** Retrieval informs quality scoring and entity resolution
5. **Its absence is measurable:** Predictions without temporal depth are provably worse

9. The Singularity Threshold

There exists a point—unique to each domain—where the accreted dataset crosses a threshold of interconnectedness and temporal depth that makes it **irreplaceable**. Not expensive to replicate. Not time-consuming. Structurally irreplaceable.

For credit reporting, this threshold was crossed decades ago. Equifax, Experian, and TransUnion have accreted payment history continuously since the 1970s. No startup, regardless of capitalization, can fabricate fifty years of consumer payment history. Their data moat is a singularity: infinite gravitational pull on the \$15 trillion lending industry because the alternative is lending without history—which is lending without risk management.

For enterprise intelligence, the threshold is closer: one to three years of continuously accreted, cross-source, entity-resolved, temporally aware data typi-

cally crosses from “useful” to “irreplaceable.” The exact point depends on the domain’s rate of change (faster-changing domains reach singularity sooner because historical patterns are more valuable) and the number of correlated sources.

9.1 The Compounding Flywheel

Once built, the accretion engine creates a self-reinforcing cycle:

*More data accreted → better entity resolution
→ richer profiles → more accurate quality
scoring → higher-value retrievals → more user
reliance → more data contributed → more data
accreted → ...*

Each revolution widens the gap between the system and any competitor that has not started its own accretion.

10. Case Study: Credit Unions vs. Bureau Scoring

The singularity thesis can be tested with a concrete case: can a cooperative network of credit unions build a data asset that rivals—or exceeds—FICO?

10.1 FICO’s Dataset

FICO scores are trained on tradeline-level data from the three major bureaus. The aggregate dataset is formidable:

Metric	Value
US consumers with scored files	282 M
Open tradelines (credit cards, auto, mortgage, etc.)	~3.4 B
Annual new originations	~550 M
Annual tradeline-level charge-offs	~18 M
Cumulative defaults observed (est. 1989–2025)	~420 M
Features per tradeline	12–15

Table 4: Estimated scale of the bureau dataset underlying FICO scoring models.

Eighteen million defaults per year across 3.4 billion tradelines provides extraordinary statistical power. This is FICO’s singularity: a label volume so large and a history so long that no new entrant can replicate it on the bureau’s own axis.

10.2 The Credit Union System

The US credit union system comprises approximately 4,500 institutions serving ~140 million members, holding ~\$1.6 trillion in outstanding loans.

Unlike bureau tradelines—which capture only account status, balance, and payment history—a credit union’s core system observes the *complete financial life* of its members:

Signal	Bureau	Credit Union
Payment history (monthly buckets)	✓	✓
Balance / utilization	✓	✓
Cross-lender visibility	✓	—
Deposit account behavior	—	✓
Income cadence & stability	—	✓
Spending category patterns	—	✓
Savings rate / liquidity buffer	—	✓
Overdraft / NSF frequency	—	✓
Digital engagement (app, channel)	—	✓
ACH / bill pay patterns	—	✓
Cash flow volatility	—	✓

Table 5: Signal comparison: bureau tradeline data vs. credit union relationship data.

The credit union sees 50–100 behavioral features per member versus the bureau’s 12–15 tradeline attributes.

10.3 Effective Information Density

In supervised learning, model quality scales with both label volume (N) and feature richness (S). Doubling feature count with high-quality, independent signals can be equivalent to 4–10× more training data. We define *effective information* as:

$$I_{\text{eff}} = N_{\text{defaults}} \times S_{\text{features}} \quad (15)$$

	Annual Defaults	Features per Obs.	I_{eff} (per year)
FICO (all bureaus)	18 M	15	270 M
FICO (CU-relevant segment)	~2.5 M	15	37.5 M
CU system (4,500 CUs)	~500 K	80	40 M
CU network (500 CUs)	~60 K	80	4.8 M

Table 6: Effective information comparison: FICO vs. credit union cooperative models.

When restricted to the CU-member-relevant segment, the full credit union system already matches FICO’s effective annual information production (~ 40 M vs. ~ 37.5 M). The gap is not $300\times$. It is approximately $1\times$ —at **year one**.

10.4 Accretion Timeline to Parity

The timeline depends on cooperative scale:

Scale	Members	FICO Parity	Superiority
1 CU (small)	10 K	Never (insufficient N)	Never
1 CU (large)	200 K	~ 36 –48 months	~ 60 months
50 CUs (CUSO)	1.5 M	~ 24 months	~ 36 months
500 CUs (network)	15 M	~ 18 months	~ 24 months
4,500 CUs (system)	140 M	~ 12 months	~ 18 months

Table 7: Projected timeline to FICO scoring parity and superiority, by cooperative scale.

Critical to this timeline is the availability of *behavioral proxy outcomes*: overdraft events (~ 500 K/year per million members), income disruptions (~ 50 K), and savings-buffer depletions (~ 100 K). These serve as leading indicators observable months before hard default, effectively accelerating label accumulation by an order of magnitude.

10.5 Where Credit Unions Win First

Parity is not uniform across segments. Credit unions achieve superiority earliest in populations where FICO is weakest:

- **Thin-file / credit invisible**: 45–50% FICO AUC vs. 65–70% CU model AUC (deposit data fills the gap)
- **Young adults** (<5 yr history): 58–62% vs. 68–72% (income/spending compensates for short credit history)
- **Recent immigrants**: 45–50% vs. 65–70% (no US bureau file; CU sees paycheck and spending)
- **Post-bankruptcy recovery**: 55–60% vs. 68–72% (bureau sees the scar; CU sees the behavioral recovery)

For steady-state prime borrowers with deep bureau files, FICO retains an advantage due to cross-lender visibility. The CU model surpasses FICO across the full membership at 36–60 months of accretion, depending on cooperative scale.

10.6 Regulatory Viability

This approach operates within existing regulatory frameworks. Under FCRA, institutions may use their own account data for credit decisions about their own members without triggering Consumer Reporting Agency obligations. Under GLBA, member NPI used for underwriting falls within “everyday business purposes.” NCUA and OCC guidance permits internal credit scoring models subject to validation, anti-lending analysis, and adverse-action-notice capability. SOC 2 Type II compliance governs the accretion infrastructure itself: encryption, access control, purpose-constrained retrieval, and immutable audit logging—precisely the architectural properties described in Section 11.

11. Architectural Requirements

The architecture required to execute this thesis has seven non-negotiable properties:

1. **Consent-bounded ingestion**: Every source registered with authorization, consent, sensitivity, purpose constraints, and retention period
2. **Deterministic policy engine**: Auditable ALLOW/REDACT/QUARANTINE/DELETE/BLOCK decisions; no ML in the policy layer
3. **Entity resolution**: Recognition and merger of records describing the same real-world entity across sources
4. **Four-tier memory compression**: Raw $\rightarrow$ Normalized $\rightarrow$ Semantic $\rightarrow$ Symbolic, each independently queryable with independent retention
5. **Temporal awareness**: Every record timestamped, every attribute change versioned, “as-of” queries supported
6. **Immutable audit trail**: Append-only log of every ingestion, decision, retrieval, merge, deletion, and export
7. **Purpose-constrained retrieval**: Queries declare purpose and actor; results filtered by source consent

The following pseudocode formalizes the three core operations. The sensitive-identifier set $\mathcal{S}_{\text{national_id}}$ is jurisdiction-configurable:

$\mathcal{S}_{\text{national_id}} = \{\text{SSN (US), ITIN (US), AADHAAR (IN), BSN (NL), STEUER_ID (DE), FINANCIAL_ACCT, \dots}\}$

Algorithm 1 INGEST — Accretion Disk Pipeline**Require:** raw input d , source record s **Ensure:** record stored or rejected with audit entry

```

1: if  $s.consent\_status \neq \text{AUTHORIZED}$  then
2:   reject  $d$ ;  $\log(\text{BLOCK}, s, \text{"consent invalid"})$ 
3:   return
4: end if
5:  $fields \leftarrow \text{PARSE}(d, s.format)$ 
6:  $pii\_tags \leftarrow \text{CLASSIFYPII}(fields) \{SSN, \text{DOB}, \text{email}, \dots\}$ 
7:  $decision \leftarrow \text{EVALUATEPOLICY}(fields, pii\_tags, s)$ 

8: if  $decision = \text{BLOCK}$  or  $decision = \text{DELETE}$  then
9:    $\log(decision, s, pii\_tags)$ 
10:  return
11: end if
12: if  $decision = \text{REDACT}$  then
13:    $fields \leftarrow \text{REDACT}(fields, pii\_tags)$ 
14: end if
15:  $canonical \leftarrow \text{NORMALIZE}(fields) \{\text{Tier 2}\}$ 
16:  $entity \leftarrow \text{RESOLVEENTITY}(canonical)$ 
17:  $semantic \leftarrow \text{SUMMARIZE}(canonical) \{\text{Tier 3}\}$ 
18:  $symbolic \leftarrow \text{EXTRACTKEYWORDS}(semantic) \{\text{Tier 4}\}$ 
19:  $vector \leftarrow \text{EMBED}(canonical)$ 
20:  $\text{STORE}(entity, [fields, canonical, semantic, symbolic, vector], s, t_{\text{now}})$ 

21:  $\log(\text{ALLOW}, s, entity, |fields|)$ 

```

Algorithm 2 EVALUATEPOLICY — Deterministic Policy Engine**Require:** fields F , PII tags P , source s **Ensure:** decision $\in \{\text{ALLOW}, \text{REDACT}, \text{QUARANTINE}, \text{BLOCK}, \text{DELETE}\}$

```

1: if  $s.sensitivity = \text{CRITICAL}$  and  $s.consent\_scope \not\supseteq \text{STORE}$  then
2:   return  $\text{BLOCK}$ 
3: end if
4: if  $\exists p \in P : p.category \in \mathcal{S}_{\text{national\_id}}$  then
5:    $qual \leftarrow (\text{UNDERWRITING} \in s.purposes) \vee (\text{COMPLIANCE} \in s.purposes)$ 
6:   if  $\neg qual$  then
7:     return  $\text{REDACT}$ 
8:   end if
9: end if
10: if  $s.retention\_days > 0$  and  $\text{AGE}(F) > s.retention\_days$  then
11:   return  $\text{DELETE}$ 
12: end if
13: if  $|P| > \theta_{\text{quarantine}}$  then
14:   return  $\text{QUARANTINE}$  {high PII density — human review}
15: end if
16: return  $\text{ALLOW}$ 

```

Algorithm 3 RETRIEVE — Purpose-Constrained Query**Require:** query q , declared purpose ρ , actor a **Ensure:** result set R filtered by consent

```

1:  $candidates \leftarrow \text{SEMANTICSEARCH}(\text{EMBED}(q), k)$ 

2: for each record  $r \in candidates$  do
3:    $s \leftarrow r.source$ 
4:   if  $\rho \notin s.allowed\_purposes$  then
5:     remove  $r$  from candidates
6:   else if  $s.consent\_expiry < t_{\text{now}}$  then
7:     remove  $r$  from candidates
8:   else
9:      $r.score \leftarrow r.similarity \times \text{FRESHNESSDECAY}(r.t, t_{\text{now}})$ 
10:  end if
11: end for
12:  $R \leftarrow \text{RANKBYScore}(candidates)$ 
13:  $\log(\text{RETRIEVE}, a, \rho, |R|, t_{\text{now}})$ 
14: return  $R$ 

```

A reference prototype demonstrating these

operations—the Data Accretion Engine (Module 07 of the Black Hole Lab simulation platform)—models the consent-bounded ingestion flow, five-decision policy engine, PII classification, tiered memory compression, and append-only audit logging. Production deployment would require hardened infrastructure, real data connectors, and validated ML pipelines, but the architectural pattern is sound and implementable with current technology.

11.1 PRE vs. PER: The Ingestion Boundary

The ingestion pipeline is divided into two critical phases, distinguished by their timing and failure modes:

PRE-ingestion checks (lines 1–3 of Algorithm 1) are synchronous, deterministic, and *blocking*:

- Consent gate: Is the source authorized? Is consent valid? Is the source enabled?
- These checks happen *before* any parsing or classification
- Failure results in immediate rejection (BLOCK decision)
- No external API calls; all checks are local rule evaluation
- Latency: $O(1)$ — typically < 10 ms

PER-ingestion processing (lines 4–11 of Algorithm 1) are asynchronous, potentially slow, and *non-blocking*:

- PII classification: Detect sensitive data categories (SSN, email, medical records, ...)
- Policy evaluation: Apply jurisdiction-specific rules to the classified data
- Redaction: Mask sensitive fields if policy permits storage with redaction
- Normalization, entity resolution, semantic compression, embedding generation
- These steps happen *during* ingestion, after the consent gate
- Failure in any PER step does not reject the ingestion; it degrades gracefully (e.g., missing embedding, failed compression)
- External API calls (LLM for classification, embedding service) are allowed
- Latency: $O(n)$ where n is document length; typically 100–500 ms for a 10 KB document

The distinction is architectural: **PRE gates are hard constraints; PER gates are soft constraints**. A source that fails the consent gate never enters the system. A source that passes the consent gate but fails PII classification still enters the system (quarantined or redacted, depending on policy). This design ensures that:

1. **Availability**: Slow external services (LLM, embedding API) cannot block ingestion
2. **Auditability**: Every document’s path through the pipeline is logged, including which PER steps succeeded or failed
3. **Graceful degradation**: Missing embeddings or failed summaries do not prevent storage; they only reduce queryability
4. **Compliance**: Hard PRE gates enforce regulatory boundaries; soft PER gates optimize for utility without sacrificing safety

12. Tropical Algebraic Structure

The policy engine and retrieval algorithm are not ad hoc conditional cascades. They instantiate a well-studied algebraic structure: the *tropical semiring* [14].

12.1 The Tropical Semiring

The tropical semiring $(\mathbb{R} \cup \{+\infty\}, \oplus, \otimes)$ replaces conventional arithmetic:

$$a \oplus b = \min(a, b) \quad (\text{tropical addition}) \quad (16)$$

$$a \otimes b = a + b \quad (\text{tropical multiplication}) \quad (17)$$

with identity elements $+\infty$ for $\oplus$ and 0 for $\otimes$. This semiring lacks additive inverses—you cannot subtract in the tropics—but it inherits associativity, commutativity of $\oplus$, and distributivity of $\otimes$ over $\oplus$.

12.2 The Policy Decision Lattice

Assign costs to the five policy decisions by restrictiveness:

Decision	Cost c_i
BLOCK	0
DELETE	1
REDACT	2
QUARANTINE	3
ALLOW	4

Table 8: Policy decisions ordered by restrictiveness.

Let each rule r_i in Algorithm 2 yield cost c_i if its guard evaluates to true, and $+\infty$ otherwise. The policy engine computes:

$$decision = \bigoplus_{i=1}^k r_i(F, P, s) = \min_i c_i \cdot \mathbf{1}[\text{guard}_i] \quad (18)$$

This is a tropical sum: the most restrictive applicable decision wins. The ALLOW at the bottom of the cascade acts as the tropical additive identity—it is returned only when all guards evaluate to $+\infty$.

This structure has a key property: **the policy lattice is a complete lattice under min**, which means:

- Rule ordering is irrelevant—the tropical sum is commutative
- Adding new rules can only make the output *more* restrictive, never less (monotonicity)
- The function is idempotent: applying the same rule twice does not change the result

12.3 Retrieval as Tropical Matrix–Vector Product

In Algorithm 3, the score computation operates in log-space:

$$\log(score_j) = \log(sim_j) + \log(\text{FRESHNESSDECAY}(\Delta t_j)) \quad (19)$$

This is tropical multiplication ($\otimes$). Ranking by maximum score across the candidate set is:

$$r^* = \bigoplus_j^{\max} score_j = \max_j score_j \quad (20)$$

using the max-plus variant of the tropical semiring. The full retrieval pipeline—embed, score, rank—is a tropical matrix–vector product:

$$R = S \otimes_{\text{max-plus}} \vec{q} \quad (21)$$

where S is the similarity–freshness matrix and $\vec{q}$ is the query embedding. This is not metaphorical: the same computation appears in shortest-path algorithms (Floyd–Warshall), Viterbi decoding, and dynamic programming over sequences.

12.4 Consent Graph Shortest Paths

Entity resolution across jurisdictions creates a *consent graph* $G = (V, E)$ where vertices are entity–source pairs and edges represent consent relationships weighted by cost (sensitivity, restriction level, jurisdictional compliance burden). Finding whether entity e_i can be linked to entity e_j under a given purpose is a shortest-path query:

$$d(e_i, e_j) = \bigoplus_{p \in \text{paths}(i,j)} \bigotimes_{(u,v) \in p} w(u, v) \quad (22)$$

which is exactly tropical matrix exponentiation: $D = W^{\otimes n}$. A consent path with cost $+\infty$ is unreachable—the entities cannot be linked under the declared purpose. This provides a formal basis for the purpose-constrained entity resolution that the paper describes informally.

12.5 Why This Matters

The tropical structure is not a curiosity. It is *why formal verification is tractable* for this system. Decision procedures over the tropical semiring are decidable: the policy engine’s behavior can be expressed as a system of tropical polynomial equations, and satisfiability of tropical polynomial systems is in NP [14]. This means:

- **Z3 can solve it:** The policy rules, expressed as tropical constraints, map directly to the quantifier-free bitvector fragment that SMT solvers handle efficiently
- **TLA⁺ can model it:** The lattice structure ensures that state transitions are monotone, eliminating a large class of liveness violations by construction

- **Counterexamples are interpretable:** A violation of a tropical invariant produces a concrete witness—a specific (F, P, s) tuple and the rule sequence that leads to the violation—making debugging auditable

The data accretion engine is not just an engineering system with post-hoc verification bolted on. Its algebraic structure *guarantees* that verification is possible, efficient, and interpretable.

13. Limitations and Threat Model

No thesis survives contact with adversaries unless its failure modes are enumerated.

13.1 Cold-Start Problem

The singularity thesis assumes accretion over time. A new system has zero data and zero temporal depth. The effective information I_{eff} at $t = 0$ is zero regardless of feature richness. Mitigation requires either (a) historical data migration from legacy systems, (b) cooperative data pooling (the CUSO model), or (c) synthetic data augmentation to bootstrap model training while real accretion proceeds.

13.2 Consent Fatigue and Attrition

Consent is the event horizon—but consent can be revoked. If member consent attrition exceeds new consent acquisition, the system’s effective mass *decreases*. This is analogous to Hawking radiation exceeding accretion: the black hole evaporates. Mitigation: continuous demonstration of value returned to members (lower rates, faster approvals, better fraud detection) to maintain consent conversion above attrition.

13.3 Data Poisoning

An adversary who can inject fabricated records into the ingestion pipeline can corrupt entity resolution, bias model training, and undermine audit integrity. The deterministic policy engine (Algorithm 2) provides a partial defense: anomalous PII density triggers QUARANTINE. But sophisticated poisoning—plausible-looking records with subtly shifted distributions—requires statistical anomaly detection at the normalization tier, which this paper does not formalize.

13.4 Regulatory Divergence

The jurisdiction-configurable identifier set $\mathcal{S}_{\text{national_id}}$ and purpose-constrained retrieval assume a consistent regulatory *grammar*: consent, purpose, retention, audit. But regulatory frameworks diverge in practice. GDPR’s “right to be forgotten” conflicts with the immutable audit trail. India’s Account Aggregator framework mandates data deletion after consent expiry, while US FCRA allows retention for “permissible purpose.” A production system must implement jurisdiction-specific retention policies that may require selective cryptographic erasure (crypto-shredding) rather than physical deletion.

Erasure has a physical floor. Landauer’s principle [1] establishes that erasing one bit of information dissipates at least $kT \ln 2$ joules of energy—approximately 3×10^{-21} J at room temperature, experimentally confirmed by Bérut et al. [15]. For a data accretion system storing 10^{12} bits per entity across 10^6 members, physical erasure of the full corpus costs on the order of 10^{-3} J—thermodynamically trivial. But the *computational* cost of locating, validating, propagating, and auditing the deletion across four memory tiers, entity resolution graphs, and downstream caches is orders of magnitude larger.

Crypto-shredding is the minimum-energy compliance strategy: encrypt each entity’s worldline under a per-entity key, and “erase” by destroying the key. The ciphertext remains in the append-only log (satisfying audit immutability), but the plaintext is irrecoverable (satisfying GDPR erasure). The energy cost reduces to erasing the key—typically 256 bits, or $\sim 5 \times 10^{-19}$ J. This is Landauer’s principle applied as an engineering constraint: **the thermodynamic cost of erasure favors crypto-shredding over physical deletion by a factor of 10^6 for a typical entity record.**

This also deepens the moat. A competitor who begins accreting today cannot reconstruct your worldlines even if they acquire the ciphertext—the keys are gone, and Landauer’s principle guarantees that the information is physically irreversible.

13.5 Feature Correlation and Overfitting

The effective information metric $I_{\text{eff}} = N \times S$ assumes features are independent. In practice, credit union behavioral signals (deposit patterns, spending categories, overdraft frequency) are correlated. The true information gain from 80 features may be equivalent to 30–40 independent features rather than the full 80. The parity timeline estimates in Table 7 should be read as optimistic bounds; empirical validation on real cooperative data is required.

13.6 Adversarial Consent

A sophisticated adversary could grant consent specifically to inject a tracking signal—a canary record designed to detect whether the system shares data beyond its stated purposes. While purpose-constrained retrieval (Algorithm 3) architecturally prevents this, implementation bugs in the consent-filtering logic would be catastrophic. Formal verification of the policy engine—not just testing—is warranted for production deployment.

13.7 Toward Formal Verification

The deterministic policy engine (Algorithm 2) is a finite-state decision function with bounded inputs. This makes it amenable to formal methods that provide *mathematical proof* of correctness, not merely empirical confidence.

1. **Model checking (TLA⁺, Alloy)**: Specify each policy decision as a state transition and each safety property as a temporal-logic invariant. Example invariant: $\Box(\text{consent_expired}(r) \implies r \notin R)$ —“no record with expired consent ever appears in a query result.” A model checker exhaustively verifies this across every reachable state of the system, not just test cases [13].
2. **SMT solving (Z3, CVC5)**: Encode the policy rules as first-order constraints and pose the negation of safety properties as a satisfiability query. If the solver returns UNSAT, the property holds for *all* inputs. This is particularly effective for the policy engine because its logic is quantifier-free bitvector arithmetic over finite enumerations.
3. **Property-based testing (QuickCheck-style)**: Generate millions of random (*fields*, *P*, *s*)

tuples and verify invariants hold. Not a formal proof, but a practical bridge: it catches the edge cases that unit tests miss (e.g., empty PII sets combined with CRITICAL sensitivity and expired retention).

4. **Dependent types (Lean 4, Idris 2)**: Implement the policy engine in a language where the type signature encodes the safety property. A function typed as $\text{RETRIEVE} : (q, \rho, a) \rightarrow \{R \mid \forall r \in R. \rho \in r.\text{purposes}\}$ literally cannot compile unless the proof holds. The policy engine’s small size (10 lines of pseudocode) makes this tractable.

The recommended verification stack for production: (1) TLA⁺ for system-level invariants across ingestion, policy, and retrieval; (2) Z3 for exhaustive input-space coverage of the policy decision function; (3) property-based testing as a continuous integration gate. The cost is modest—the policy engine is deliberately small and deterministic precisely to enable this kind of verification.

14. Conclusion: The 426-Byte Endgame

We began with 426 bytes. A single Metro 2 trade-line. A snapshot of one financial relationship at one moment in time.

426 bytes is worthless in isolation. It has no trajectory, no context, no predictive power.

But the credit bureaus understood something fundamental: accrete 426-byte records across every creditor, for every consumer, every month, for decades, and you do not get a large pile of snapshots. You get a **singularity**—a dataset so temporally deep, so interconnected, and so irreplaceable that an entire \$15 trillion lending industry orbits around it. Three companies, accreting 426 bytes at a time, built one of the most powerful data moats in economic history.

The physics is the strategy:

- The **event horizon** (consent boundary) determines what enters the system
- The **accretion disk** (pipeline) transforms raw data into structured knowledge
- The **singularity** (memory core) compresses knowledge into irreducible density

- **Hawking radiation** (retrieval) extracts compounding value from the core
- **Time dilation** (temporal depth) ensures the moat widens with every passing day

Data accretion is not a feature. It is not an optimization. It is the recognition that **information follows the same gravitational mechanics as matter**—and that the organizations which begin accreting earliest, most consistently, and most carefully will build singularities that reshape the industries around them.

*426 bytes is the seed. The moat is the accretion.
The singularity is the endgame.*

References

- [1] R. Landauer, “Irreversibility and heat generation in the computing process,” *IBM Journal of Research and Development*, vol. 5, no. 3, pp. 183–191, 1961.
- [2] J. D. Bekenstein, “Universal upper bound on the entropy-to-energy ratio for bounded systems,” *Physical Review D*, vol. 23, no. 2, pp. 287–298, 1981.
- [3] S. W. Hawking, “Particle creation by black holes,” *Communications in Mathematical Physics*, vol. 43, no. 3, pp. 199–220, 1975.
- [4] C. E. Shannon, “A mathematical theory of communication,” *Bell System Technical Journal*, vol. 27, no. 3, pp. 379–423, 1948.
- [5] E. Warren, “The economics of consumer credit,” *Connecticut Law Review*, vol. 36, pp. 1235–1275, 2004.
- [6] R. Penrose, “Gravitational collapse and space-time singularities,” *Physical Review Letters*, vol. 14, no. 3, pp. 57–59, 1965.
- [7] L. Susskind, “The world as a hologram,” *Journal of Mathematical Physics*, vol. 36, no. 11, pp. 6377–6396, 1995.
- [8] M. E. Porter, *Competitive Advantage: Creating and Sustaining Superior Performance*. New York: Free Press, 1985.
- [9] FinRegLab, “The Use of Cash-Flow Data in Underwriting Credit: Empirical Research Findings,” Research Report, 2019. Available: <https://finreglab.org/cash-flow-data>
- [10] Consumer Financial Protection Bureau, “An Updated Review of the New and Revised Data Points in HMDA,” Research Report, 2019.
- [11] National Credit Union Administration, “Quarterly Credit Union Data Summary,” Q4 2024. Available: <https://ncua.gov/analysis>
- [12] L. Braille, *Procédé pour écrire les paroles, la musique et le plain-chant au moyen de points*. Paris: Institution Royale des Jeunes Aveugles, 1829.
- [13] L. Lamport, *Specifying Systems: The TLA⁺ Language and Tools for Hardware and Software Engineers*. Boston: Addison-Wesley, 2002.
- [14] D. Maclagan and B. Sturmfels, *Introduction to Tropical Geometry*, Graduate Studies in Mathematics, vol. 161. Providence: AMS, 2015.
- [15] A. Bérut *et al.*, “Experimental verification of Landauer’s principle linking information and thermodynamics,” *Nature*, vol. 483, pp. 187–189, 2012.